

# I have a good shard key now what?

David Murphy , Mongo Master  
Lead DBA, ObjectRocket  
[@dmurphy\\_data](#) [@objectrocket](#)



# Background

- 16 yrs in databases, development, & system engineering
- Lead DBA @ ObjectRocket
- Mongo Master with a focus on sharding, chunks, and scaling mongo beyond normal means.



# Quick Sharding Introduction

## Why you might need to shard:

- Scaling Reads
- Fast Queries
- Scaling Writes
- Balancing nodes

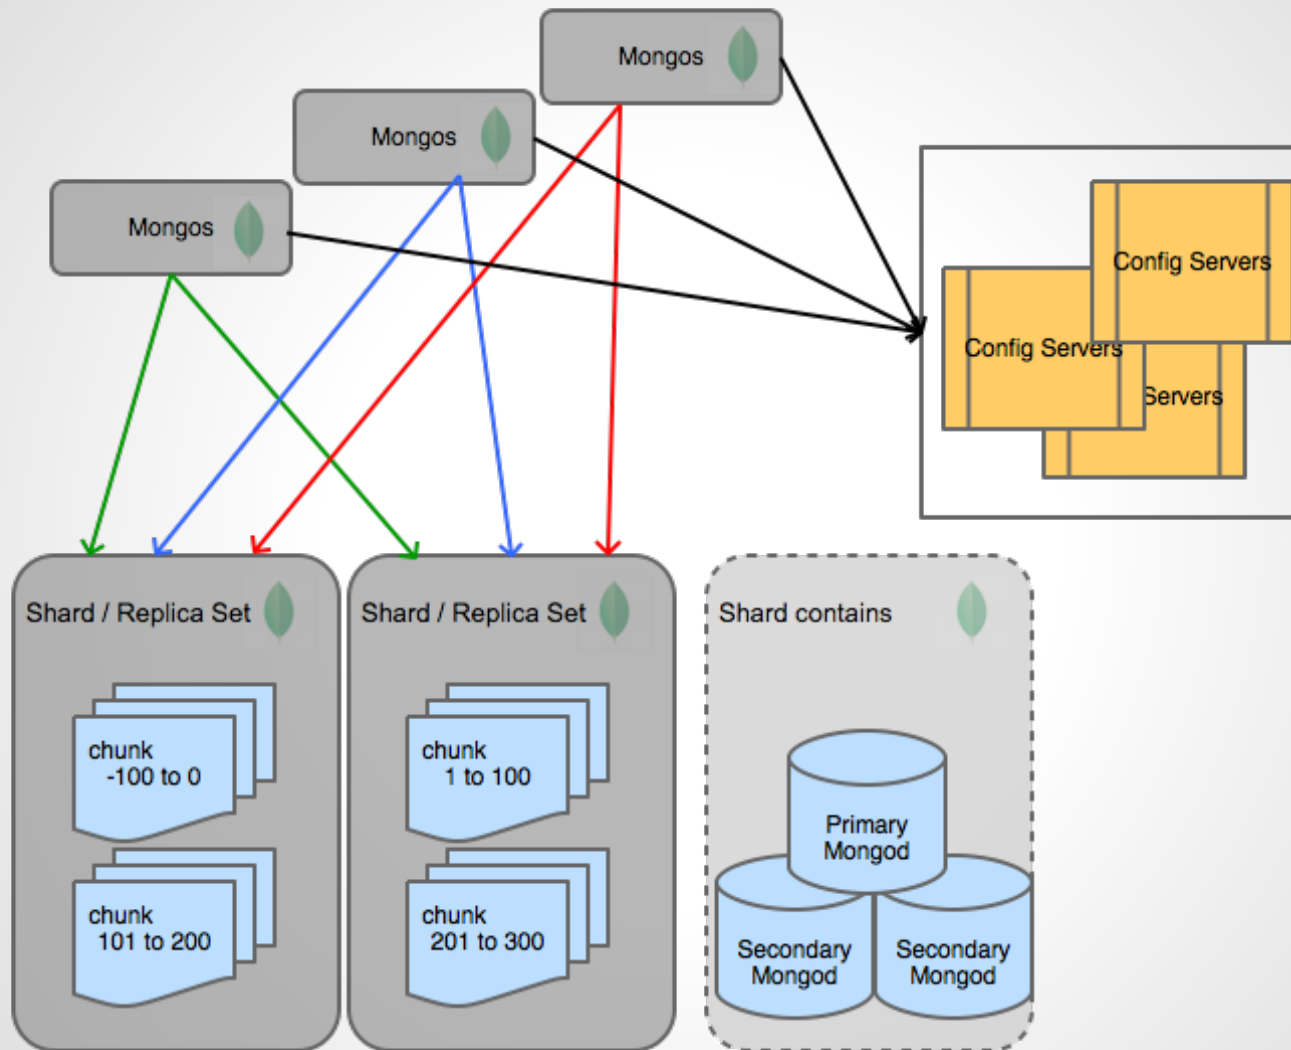
## What are things should you consider?

- Shard Key is Immutable
- Targeted vs Scatter-Gather
- Sharding Overhead
- Some commands no longer work

<http://bit.ly/1oXYDfm> - Kenny Gorman sharding talk

<http://bit.ly/ZTtDI1> - Several other sharding talks

# What does a shared cluster look like?



# The balancer is great...

*... but you can also perform some proactive checks to predict issues.*

- ☐ Understanding the reasons why not all chunks are equal
- ☐ Counting Chunks per shard without `sh.status()`
- ☐ Finding out how many `chunkMove`'s or splits occur at a time
- ☐ Getting the size of all chunks for a collection



# Ways to Count Chunks

Many functions for this:

- `sh.status()`
- `db.collection.stats()`

Simple query for counting chunks:

- `db.chunks.count({ns:"database.collection"});`

Why I prefer using Aggregation:

- Faster as you only look at one type of data
- Provides chunk count per shard
- Can programmatically use the output



# Tracking Chunks and Moves

The **config.changelog** tracks chunk changes.

The types of changes to chunks are

- **split**
- **moveChunk**
  - **moveChunk.start**
  - **moveChunk.from**
  - **moveChunk.to**
  - **moveChunk.commit**

Chunk Collection Examples

<http://bit.ly/112tXEx>

ChangeLog Examples

<http://bit.ly/1oZlyqN>



# Sizing Chunks

- dataSize command - Very expensive  
Scans all documents to be 100% accurate
- find().count() \* avgObjSize on chunk ranges  
Rough size estimate based on stats
- ChunkHunter.py (On GitHub <http://bit.ly/1yWQf9T> )  
New community script from ObjectRocket to:
  - a) copy chunks meta data to new collection
  - b) scan each chunk there with dataSize or count
  - c) output summary





# Example ChunkHunter Output

```
[root@mon2 ~]# python ChunkHunter.py -H $DBHOST -u $DBUSER -P $DBPORT -p $DBPASS \  
-d ChunkHunterTest -c ChunkHunter -O "test.foo" -m datasize
```

```
Chunk count for config.chunks was 4251
```

```
Populating Chunks from config database:
```

```
    Populating 282 of 282 documents
```

```
Chunk count for test.foo was 282
```

```
Processing Chunks for size using datasize:
```

```
    Processing 282 of 282
```

```
Summary Report
```

```
Chunk count for test.with_jumbos was 282
```

-----+-----+-----+-----+
Namespace   Count   Jumbo by Doc   Jumbo by Size
-----+-----+-----+-----+
ChunkHunterTest.ChunkHunter   282   19   93
-----+-----+-----+-----+

```
We recommend checking the output as you may have splittable chunks to improve your balance
```



# Example ChunkHunter Output

```
[root@mon2 ~]# python ChunkHunter.py -H $DBHOST -u $DBUSER -P $DBPORT -p $DBPASS \  
-d ChunkHunterTest -c ChunkHunter -O "test.foo" -m datasize
```

Chunk count for config.chunks was 4251 Total Chunks in Cluster

Populating Chunks from config database:

Populating 282 of 282 documents

Chunk count for test.foo was 282

Processing Chunks for size using datasize:

Processing 282 of 282

Summary Report

Chunk count for test.with\_jumbos was 282

Namespace	Count	Jumbo by Doc	Jumbo by Size
ChunkHunterTest.ChunkHunter	282	19	93

We recommend checking the output as you may have splittable chunks to improve your balance



# Example ChunkHunter Output

```
[root@mon2 ~]# python ChunkHunter.py -H $DBHOST -u $DBUSER -P $DBPORT -p $DBPASS \  
-d ChunkHunterTest -c ChunkHunter -O "test.foo" -m datasize
```

```
Chunk count for config.chunks was 4251
```

```
Populating Chunks from config database:
```

```
    Populating 282 of 282 documents
```

Total Chunks Populated

```
Chunk count for test.foo was 282
```

```
Processing Chunks for size using datasize:
```

```
    Processing 282 of 282
```

```
Summary Report
```

```
Chunk count for test.with_jumbos was 282
```

-----+-----+-----+-----+-----+   Namespace	Count	Jumbo by Doc	Jumbo by Size	
-----+-----+-----+-----+-----+   ChunkHunterTest.ChunkHunter	282	19	93	
-----+-----+-----+-----+-----+				

```
We recommend checking the output as you may have splittable chunks to improve your balance
```



# Example ChunkHunter Output

```
[root@mon2 ~]# python ChunkHunter.py -H $DBHOST -u $DBUSER -P $DBPORT -p $DBPASS \
-d ChunkHunterTest -c ChunkHunter -O "test.foo" -m datasize
```

```
Chunk count for config.chunks was 4251
```

```
Populating Chunks from config database:
```

```
    Populating 282 of 282 documents
```

```
Chunk count for test.foo was 282
```

```
Processing Chunks for size using datasize:
```

```
    Processing 282 of 282 Total Chunks processed
```

```
Summary Report
```

```
Chunk count for test.with_jumbos was 282
```

+-----+-----+-----+-----+
Namespace   Count   Jumbo by Doc   Jumbo by Size
+-----+-----+-----+-----+
ChunkHunterTest.ChunkHunter   282   19   93
+-----+-----+-----+-----+

```
We recommend checking the output as you may have splittable chunks to improve your balance
```



# Example ChunkHunter Output

```
[root@mon2 ~]# python ChunkHunter.py -H $DBHOST -u $DBUSER -P $DBPORT -p $DBPASS \
-d ChunkHunterTest -c ChunkHunter -O "test.foo" -m datasize
```

```
Chunk count for config.chunks was 4251
```

```
Populating Chunks from config database:
```

```
    Populating 282 of 282 documents
```

```
Chunk count for test.foo was 282
```

```
Processing Chunks for size using datasize:
```

```
    Processing 282 of 282
```

```
Summary Report
```

>250k Documents per chunk

```
Chunk count for test.with_jumbos was 282
```

Namespace	Count	Jumbo by Doc	Jumbo by Size
ChunkHunterTest.ChunkHunter	282	19	93

```
We recommend checking the output as you may have splittable chunks to improve your balance
```





# Example ChunkHunter Output

```
[root@mon2 ~]# python ChunkHunter.py -H $DBHOST -u $DBUSER -P $DBPORT -p $DBPASS \
-d ChunkHunterTest -c ChunkHunter -O "test.foo" -m datasize
```

```
Chunk count for config.chunks was 4251
```

```
Populating Chunks from config database:
```

```
    Populating 282 of 282 documents
```

```
Chunk count for test.foo was 282
```

```
Processing Chunks for size using datasize:
```

```
    Processing 282 of 282
```

```
Summary Report
```

>64M per Chunk

```
Chunk count for test.with_jumbos was 282
```

Namespace	Count	Jumbo by Doc	Jumbo by Size
ChunkHunterTest.ChunkHunter	282	19	93

```
We recommend checking the output as you may have splittable chunks to improve your balance
```



# What did we add to the doc?

```
{
  "_id" : "ChunkHunterTest.ChunkHunter-_id_MinKey",
  "ns" : "ChunkHunterTest.ChunkHunter",
  "min" : {
    "_id" : { $minKey : 1 }
  },
  "docs" : 37213,
  "shard" : "0e5302a229a01e20cf4e29ae4f352c54",
  "processed" : true,
  "max" : {
    "_id" : NumberLong("-9156596508897956698")
  },
  "jumbo" : false,
  "size" : 3.09332275390625
}
```



# Some example queries might be

- Give me all Jumbo:true chunks
- Order Jumbo chunks by documents so I can split them
- Order Jumbo chunks by size so I can split them
- Aggregate Jumbo chunks to count them per shard
- Give me all Jumbo on Shard “X”

Find all these and more at:

<http://bit.ly/1oYVGeJ>





# The Future...

Will be releasing tools to simplify:

- Manual Splits
- Moving of Chunks

These are not replacements for the balancer but ways to help it stay on track if things deviate. As opposed to waiting for an incident and fixing it in a panic.



# Contact

@dmurphy\_data  
@objectrocket  
david@objectrocket.com  
<https://www.objectrocket.com>

**WE ARE HIRING! (DBA,DEVOPS, and more)**  
<https://www.objectrocket.com/careers>

## Resources:

ChunkHunter.py <http://bit.ly/1yWQf9T>  
Chunk Hunter Example Scripts [bit.ly/1oYVGeJ](http://bit.ly/1oYVGeJ)  
Chunk Collection Queries <http://bit.ly/112tXEx>  
Changelog Queries <http://bit.ly/1oZlyqN>

## Presentations:

Kenny Gorman Sharding - [bit.ly/1oXYDfm](http://bit.ly/1oXYDfm)  
Other Sharding Links - [bit.ly/ZTtDI1](http://bit.ly/ZTtDI1)

